

NLP Milestone

30093813

ACM Reference Format:

30093813. 2023. NLP Milestone. 1, 1 (March 2023), 3 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

INTRODUCTION

This project will attempt to address the problems resulting from the necessity for manual linking in educational material. In particular, educational material often requires large amounts of repetitive linking to related explanatory material as it uses domain-specific terminology. This project will explore the NLP techniques that can be used to enable the automatic identification of the spans of a document which should be linked, and the pairing of such spans with their appropriate links.

SOLUTION

The work on the project so far has consisted mainly of data preparation, with some minor initial exploration done into model training with SpaCy.

The dataset chosen was the `scikit-learn` documentation (GitHub source: <https://github.com/scikit-learn/scikit-learn/tree/main/doc>, commit 449940985c903f77678c0627cbc7a6267c3a54f9). To extract the link data from the documents, I wrote a tool which uses Pandoc to assign UUIDs to each link in each document, wrap the link content with a special marker for later extraction, store the UUID-to-link mapping in a JSON file, and convert the HTML to plain text. Pandoc was chosen to make it easier to use the tool on different datasets with different documentation formats.

After that, Python is used for further processing of the links. The dataset consists of 996 documents with a total length of 1100122 words (according to the `wc` command line utility). Unprocessed, the dataset consists of a total of 17982 unique links. Multiple processing steps were used to account for different link forms: the links were lowercased, "http" forms of existing "https" links were converted to "https", Python's `urllib` was used for normalization, and relative links were normalized. Following this processing, the dataset consisted of 17164 unique links.

The majority of links have very few occurrences; 1 illustrates the exponentially decreasing relationship between the minimum occurrence count and the number of classes meeting that minimum. More than half the links in the dataset have only one example, and only 474 out of the ~17000 links have more than 10 examples.

To help reduce the number of classes, we can focus on the links with many examples; fortunately this is suitable for the application, because links with many examples are more likely to be used again.

There is also the issue of splitting the data into training and testing data. For this task I wrote a script which provides a roughly balanced split of the documents in the dataset, in terms of their example counts for each class. The script accepts a test size and a minimum example threshold, and performs a search for a balanced split. The split class statistics

Author's address: 30093813.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2023 Association for Computing Machinery.

Manuscript submitted to ACM

Manuscript submitted to ACM




Fig. 1. Distribution of minimum example count

are summarized in 1. A minimum example threshold will be chosen based on model results, but these initial findings indicate that 15 may be a good threshold.

Table 1. Counts of classes in training and test sets for different thresholds

Minimum examples	# Training set classes	# Testing set classes	# Classes with no test examples
5	1152	1029	123
10	474	466	8
15	274	273	1
20	181	180	1
25	123	122	1
30	104	104	0

After the data is split into training and testing sets, another script extracts the link locations from the plain text files and creates SpaCy Doc's annotated with the link locations and their link labels. These Doc's are stored in SpaCy's binary data format so that a SpaCy pipeline can be trained on the data.

Thus far, a single pipeline has been trained on the minimum 15 threshold dataset with a tok2vec component and a spancat (span categorizer) component predicting links, with the best model reaching an F-score of ~0.8. This score may not reflect the actual utility of the model, as the extraneous links described in Challenges may be boosting the

score. It was not possible to train the pipeline with the default configuration on unpaid Google Colab due to limited memory, but I was able to train it on my laptop on CPU. This pipeline was an initial experiment to confirm the data was in an appropriate format for training. Over the following month I will be further researching the training process in SpaCy and evaluating different models.

CHALLENGES

The documentation for `scikit-learn` is in RST format, and it has a build tool in place which allows special link types. This makes it so that the full hyperlinks cannot be parsed directly from the RST source, so instead I have taken the built HTML source and parsed the links from there. This has the downside that extraneous hyperlinks, such as those in the navigation bar, may need to be excluded as a special case. This special case will be explored based on the model results. If these hyperlinks need to be removed, it will be possible by limiting the parsed HTML to only those within the tag representing the main content of the webpage.

Another unanticipated challenge came with splitting the documents into train and test data such that the classes were balanced. The combination of a large number of classes, a large imbalance in the number of examples for the different classes, and an unequal representation of the classes across documents made it challenging to find a balance, and I couldn't find an existing tool for the task. I was able to write a search for this, but it doesn't always find a perfect balance. If the split ends up being unsuitable, the search can be applied to split by paragraph instead of document, or possibly by sentence, depending on the amount of context required by the model.