

# NLP Project Report: Corpus-Specific Automatic Hyperlinking

30093813

## ACM Reference Format:

30093813. 2023. NLP Project Report: Corpus-Specific Automatic Hyperlinking. 1, 1 (April 2023), 12 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

## INTRODUCTION

### Problem definition

This project will attempt to address the problems resulting from the necessity for manual linking in educational material. In particular, educational material often requires large amounts of repetitive linking to related explanatory material as it uses domain-specific terminology. This results in at least one of two sets of problems:

- (1) The authors of the educational material under-link, meaning they use domain-specific terminology without linking to its corresponding explanatory material, resulting in decreased effectiveness of the educational material as its consumers now face a higher barrier-to-entry.
- (2) The authors of the educational material manually link, resulting in decreased efficiency during authoring, a higher chance of error in linking, and an increased cost of maintenance as the educational material evolves and links change.

An example corpus of educational material which suffers from these problems is the `scikit-learn` [9] [documentation](#). The documentation contains a wealth of knowledge relating not only to the library itself but machine learning as a whole, however for a reader who is new to the domains of machine learning and statistics it can be quite opaque as it consistently uses domain specific terminology. For example, the first sentence in the user guide reads:

“The following are a set of methods intended for regression in which the target value is expected to be a linear combination of the features.”

There isn’t anything inherently wrong with this as the documentation assumes working knowledge of the domain, however it has potential for improvement. For example, the glossary already contains definitions for the terms “target” [14] and “feature” [12], and even has an example of “feature” being used and linked in context [13]. However a reader who is new to the domain who stumbles across such articles will likely be confused, and educational material would ideally provide an obvious path for resolving that confusion.

---

Author’s address: 30093813.

---

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

© 2023 Association for Computing Machinery.

Manuscript submitted to ACM

Manuscript submitted to ACM

## Approach

This project attempts to address the root of the problem: it is easier to not link than to link. The project explores the NLP techniques that can be used to enable the automatic identification of the spans of a document which should be linked, and the pairing of such spans with their appropriate links.

The project uses the `scikit-learn` [9] documentation as a dataset. It parses the links from the HTML using Pandoc [8], processes the data using Python, and trains a span classification pipeline using SpaCy [4].

The project is in essence a supervised classification task, and as such it will be evaluated using classification metrics such as those listed in the `scikit-learn` user guide [11].

## Related Work

This project falls under the category of text based entity linking [19]. More specifically, it is essentially a system to support corpus-specific “wikification” of any corpus with parse-able links (as opposed to being limited to linking to Wikipedia). Existing approaches to text based entity linking typically follow a two step process of first identifying entities (a.k.a. keyword extraction) then matching them to their links (a.k.a word sense disambiguation), and focus on using Wikipedia data [1, 6, 7]. Existing approaches also make use of the additional metadata provided by Wikipedia datasets, such as the entity categories and classes, to help with the common problem of disambiguation. Existing approaches also make heavy use of the existing links within the corpus to extract the context and “surface forms” (different terms for the same entity) of the entities. The most difficult problem appears to be that of disambiguation. Yin et al. seems to represent the state-of-the-art in entity linking and uses the newer technique of transformer models with BERT to learn the context of the entities in a document. With SpaCy’s recent support for training transformer models in pipelines, I explore a similar approach in this project.

## MATERIALS AND METHODS

### Dataset and Link Extraction

The dataset chosen was the `scikit-learn` documentation (GitHub source: <https://github.com/scikit-learn/scikit-learn/tree/main/doc>, commit 449940985c903f77678c0627cbc7a6267c3a54f9). To extract the link data from the documents, I wrote a tool which uses Pandoc [8] to assign UUIDs to each link in each document, wrap the link content with a special marker for later extraction, store the UUID-to-link mapping in a JSON file, and convert the HTML to plain text. Pandoc was chosen to make it easier to use the tool on different datasets with different documentation formats. The tool also has the option to narrow the link extraction to a particular element within the HTML; this was used to only extract the main content of the pages, while excluding the other HTML like the navigation bar. This had the effect of greatly increasing link balance and relevance, as the navigation elements were present on the majority of pages, and they contained the same links on every page.

### Link Normalization

After the link data and raw text is extracted, Python is used for further processing of the links. The dataset consists of 996 documents with a total length of 1100122 words (according to the `wc` command line utility). Unprocessed, the dataset consists of a total of 16905 unique links. Multiple processing steps were used to account for different link forms: the links were lowercased, “http” forms of existing “https” links were converted to “https”, Python’s `urllib` was used

for normalization, and relative links were normalized. Following this processing, the dataset consisted of 16162 unique links.

### Link Occurrence Filtering

Links occurrences were further filtered by removing cases where the occurrence’s user readable text did not contain at least 3 alphanumeric characters. This was implemented to prevent the model from trying to learn from uninformative text which does not have any special relation to the particular link it corresponds to. For example, the corpus contains many occurrences of “¶”, which is used generally to link to headers, so that same span is associated with many different links. Furthermore, these non-alphanumeric spans are most likely generated text as opposed to user authored text, so there isn’t any need to automatically link these spans.

Following this filtering, the dataset contained 9715 unique links.

### Generating Training Examples

The dataset has the property that the majority of the text is a negative sample (it does not contain any links). To deal with this, training examples were generated as follows:

- (1) A sliding window with a width of three sentences was run over each document. If the window contains any link spans, that window is added as a training example. Sentences were predicted by SpaCy’s [Sentencizer](#). Other sentence splitting strategies provided by SpaCy don’t seem to work as well, likely due to the unique formatting of the corpus (for example it includes code blocks).
- (2) Cases where the links cross sentence boundaries and are not included in a training example are given a context of 10 words on either side, and this is included as a training example.

This process results in 11315 training examples.

### Reducing Class Count

After the above pre-processing, the dataset still contains far too many classes for classification.

The majority of links have very few occurrences; Figure 1 illustrates the exponentially decreasing relationship between the minimum occurrence count and the number of classes meeting that minimum. Over 80% of the links in the dataset have only one example, and only 157 links have more than 10 examples. Table 1 summarizes the number of classes for a few milestone thresholds.

Table 1. Counts of classes for different thresholds

Minimum examples	Number of classes
5	453
10	157
15	89
20	49
25	30
30	23

To help reduce the number of classes, we can focus on the links with many examples; fortunately this is suitable for the application, because links with many examples are more likely to be used again. The majority of experiments used a




Fig. 1. Distribution of minimum example count

threshold of 15 as this threshold provides a good balance between the number of training samples per link and the number of classes to predict, making for a reasonable task for the model. A threshold of 15 also allows for a reasonable number of links, improving the model's overall utility.

### Train/Test Splitting

To split the data into training and testing sets, the main challenge is in keeping the classes balanced, as a single training example can contain multiple class counts. Originally I wrote a custom script to search for a balanced split. This script also operated on entire documents, making it harder to balance classes as the classes came in large, fixed clusters. However, if we view the problem as a multi-label classification problem for the purposes of splitting (even though this is not accurate for the actual final task of the model), we can utilize the pre-existing library scikit-multilearn [17] to do the splitting. In particular, the library provides an implementation of a multi-label stratification algorithm [15, 18]. To utilize the algorithm, one can one-hot encode the label space with a  $n \times m$  matrix, where  $n$  is the number of training examples, and  $m$  is the number of classes. This is not a perfect representation as a single training example may contain multiple instances of the same link, however the algorithm ends up working very well for the task, providing a balanced split with each class containing roughly proportionate amounts in the training and testing sets. A test size of 33% is used for most experiments.

Following the training and testing split, under-sampling is performed to further balance the class counts, as some classes are far more prevalent than others and we don't want the model to become biased towards these links. I tried to use imbalanced-learn's [5] `RandomUnderSampler`, however the library does not support multi-label datasets, and while there are ways to work around this, they ultimately don't provide a great balance of classes.

I found that manual under-sampling using the following algorithm provides a good balance:

- (1) Keep a running count of the number of occurrences for each class, and find the minimum number of occurrences for a class
- (2) For each document, count the number of occurrences of each class in the document and remove the document if subtracting those occurrences keeps all the classes above the minimum

For this particular dataset, this algorithm provided a perfect balance (although that may not be the case for other datasets). With a minimum threshold of 15 and a test size of 33%, all 89 classes in the training set contained 10 examples each, and all classes in the testing set contained 5 examples each, with both the training and testing sets having the same sets of classes. The dataset then consists of 786 training samples and 406 testing samples, 55740 total words (5575 unique), with spans containing 1 to 3 tokens with distribution: 1 (45.98%), 2 (27.23%), 3 (26.79%). The 10 most common span tokens are: ".", "pipeline", "fit", "gridsearchcv", "linear\_model", "pca", "sgdclassifier", "ensemble", "sklearn.model\_selection", "sklearn.ensemble".

## Model Training

SpaCy [4] is used for building a classifier. All training and testing data is stored in SpaCy's binary data format using `DocBin`, and SpaCy's built-in training process is used for all training. The full model configuration can be found on [GitHub](#). The pipeline consists of SpaCy's default tokenizer followed by two components: a pre-trained DistilBERT transformer [10] from the Hugging Face [3] library, followed by SpaCy's `SpanCategorizer` component.

The DistilBERT transformer was chosen to allow for the use of span context during classification, in the hopes of creating a more sophisticated classification strategy. It is also well suited for this application due to its reduced size, and its potential for speed improvements on CPUs with techniques such as those described in Shen et al.. The span categorizer is set up to categorize the spans generated during the data processing above, and it is connected to the transformer. Both components are trained together, with the span categorizer using the transformer's outputs as features. All training was performed on unpaid Google Colab GPUs.

The full training command used is:

```
python -m spacy train config.cfg --output ./output --paths.train ./train.spacy --paths.dev ./test.spacy --gpu-id 0
```

Minimal hyper-parameter tuning was performed due to time constraints and because SpaCy's defaults provided good results. One difficulty encountered was running out of GPU memory when training on the full documents, and some tuning of batch sizes was required to reduce memory usage. However after switching to smaller training examples, the memory requirements were much smaller and no longer an issue.

## RESULTS

Multiple iterations of the model were trained as the data preparation progressed. For the results, I will focus on the results of the model with the final (and best) data preparation (described above). The overall classification scores are shown below:

```

261 import json
262 with open('output/model-best9/meta.json') as f:
263     data = json.load(f)['performance']
264
265 print('Precision:', data['spans_sc_p'])
266 print('Recall:    ', data['spans_sc_r'])
267 print('F1-score: ', data['spans_sc_f'])
268
269 Precision: 0.8923766816
270 Recall:    0.8943820225
271 F1-score:  0.8933782267
272
273     Next are the scores for each class.
274
275 import pandas as pd
276 import os
277 import re
278 data = pd.read_csv(os.path.join('output', 'model-best9', 'results.csv'), index_col=0)
279 mlen = 60
280 def trunc_link(link):
281     prefix = '...' if len(link) > mlen else ''
282     return prefix + re.sub('_', ' ', link)[-min(len(link), mlen):]
283 data.index = data.index.map(trunc_link)
284
285

```

Many of the classes end up with perfect scores. The proportion of such classes is given:

```

287 f'{100 * data[data["f"] == 1].shape[0] / data.shape[0]:.2f}%'
288
289 44.94%

```

Table 2 contains the scores for the classes which did not receive perfect scores. Note that the “support” column contains the number of samples tested, and “p”, “r”, and “f” are precision, recall, and F1-score, respectively.

```

293 df = data[(data['f'] != 1)]
294 df
295

```

As an example of the model’s results, we can use SpaCy’s Displacy tool to render some predictions for the link “glossary.html#term-fit”. In the following images, “Reference” shows a sample from the testing set with its links bolded, and “Predicted” shows the model’s predictions bolded. Figure 2 shows a case containing multiple unique links and multiple occurrences of the same link, and the model fails to identify any links. Figure 3 shows another test case where the model made a prediction for the same span of text as the reference.

## DISCUSSION AND CONCLUSION

The results indicate that the system has some potential, however there are some biases in the dataset that need to be considered when evaluating the overall utility of the model. In particular, the dataset is unbalanced with far more examples of links for unique names, such as contributor names or function and class names, as opposed to terms such as “target” and “fit”. This is due largely to the inclusion of all pages of the documentation in the dataset, including changelogs (where contributors are frequently mentioned) and API documentation (where function and class names are frequently linked). It is also more common for function and class names to be linked in the documentation than

Table 2. Results without perfect F1-score

	p	r	f	support
glossary.html#term-decision function	0.5	0.2	0.285714	5
<a href="https://twitter.com/ogrisel">https://twitter.com/ogrisel</a>	0.5	0.4	0.444444	5
...ction.gridsearchcv.html#sklearn.model selection.gridsearchcv	0.5	0.4	0.444444	5
...ted/sklearn.pipeline.pipeline.html#sklearn.pipeline.pipeline	0.666667	0.4	0.5	5
...ction.gridsearchcv.html#sklearn.model selection.gridsearchcv	0.428571	0.6	0.5	5
glossary.html#term-predict	1	0.4	0.571429	5
#term-fit	0.6	0.6	0.6	5
glossary.html#term-predict proba	0.6	0.6	0.6	5
...ted/sklearn.decomposition.pca.html#sklearn.decomposition.pca	0.75	0.6	0.666667	5
glossary.html#term-fit	0.75	0.6	0.666667	5
...ted/sklearn.decomposition.pca.html#sklearn.decomposition.pca	0.571429	0.8	0.666667	5
...learn.linear model.ridgecv.html#sklearn.linear model.ridgecv	1	0.6	0.75	5
...lectfrommodel.html#sklearn.feature selection.selectfrommodel	1	0.6	0.75	5
<a href="https://sites.google.com/site/peterprettenhofer/">https://sites.google.com/site/peterprettenhofer/</a>	0.625	1	0.769231	5
<a href="http://www.montefiore.ulg.ac.be/~glouppe/">http://www.montefiore.ulg.ac.be/~glouppe/</a>	0.8	0.8	0.8	5
...arn.pipeline.featureunion.html#sklearn.pipeline.featureunion	0.8	0.8	0.8	5
...izedsearchcv.html#sklearn.model selection.randomizedsearchcv	0.8	0.8	0.8	5
... model.sgdclassifier.html#sklearn.linear model.sgdclassifier	0.8	0.8	0.8	5
generated/sklearn.svm.linearsvc.html#sklearn.svm.linearsvc	0.714286	1	0.833333	5
<a href="https://gael-varoquaux.info">https://gael-varoquaux.info</a>	0.714286	1	0.833333	5
<a href="https://github.com/ogrisel">https://github.com/ogrisel</a>	0.714286	1	0.833333	5
...ose.columntransformer.html#sklearn.compose.columntransformer	1	0.8	0.888889	5
<a href="https://github.com/thomasjpfan">https://github.com/thomasjpfan</a>	1	0.8	0.888889	5
<a href="http://www.mblondel.org">http://www.mblondel.org</a>	1	0.8	0.888889	5
...learn.impute.simpleimputer.html#sklearn.impute.simpleimputer	1	0.8	0.888889	5
...n.naive bayes.gaussiannb.html#sklearn.naive bayes.gaussiannb	1	0.8	0.888889	5
classes.html#module-sklearn.metrics.pairwise	1	0.8	0.888889	5
<a href="https://github.com/micky774">https://github.com/micky774</a>	1	0.8	0.888889	5
...gclassifier.html#sklearn.ensemble.gradientboostingclassifier	1	0.8	0.888889	5
...featurehasher.html#sklearn.feature extraction.featurehasher	1	0.8	0.888889	5
...orestclassifier.html#sklearn.ensemble.randomforestclassifier	1	0.8	0.888889	5
...tml#sklearn.discriminant analysis.lineariscriminantanalysis	1	0.8	0.888889	5
...impute.iterativeimputer.html#sklearn.impute.iterativeimputer	1	0.8	0.888889	5
<a href="https://amueller.github.io/">https://amueller.github.io/</a>	0.833333	1	0.909091	5
...ted/sklearn.pipeline.pipeline.html#sklearn.pipeline.pipeline	0.833333	1	0.909091	5
<a href="https://github.com/qinhanmin2014">https://github.com/qinhanmin2014</a>	0.833333	1	0.909091	5
... model.sgdclassifier.html#sklearn.linear model.sgdclassifier	0.833333	1	0.909091	5
<a href="http://alexandre.gramfort.net">http://alexandre.gramfort.net</a>	0.833333	1	0.909091	5
...ssing.onehotencoder.html#sklearn.preprocessing.onehotencoder	0.833333	1	0.909091	5
<a href="http://fa.bianp.net">http://fa.bianp.net</a>	0.833333	1	0.909091	5
...localoutlierfactor.html#sklearn.neighbors.localoutlierfactor	0.833333	1	0.909091	5
...essor.html#sklearn.gaussian process.gaussianprocessregressor	0.833333	1	0.909091	5
<a href="https://github.com/larsmans">https://github.com/larsmans</a>	0.833333	1	0.909091	5
...stimator.html#sklearn.utils.estimator checks.check estimator	0.833333	1	0.909091	5
generated/sklearn.svm.svr.html#sklearn.svm.svr	0.833333	1	0.909091	5
...nerated/sklearn.svm.oneclasssvm.html#sklearn.svm.oneclasssvm	0.833333	1	0.909091	5
...isticregression.html#sklearn.linear model.logisticregression	0.833333	1	0.909091	5
<a href="https://github.com/lorentzenchr">https://github.com/lorentzenchr</a>	0.833333	1	0.909091	5
...letransformer.html#sklearn.preprocessing.quantiletransformer	0.833333	1	0.909091	5

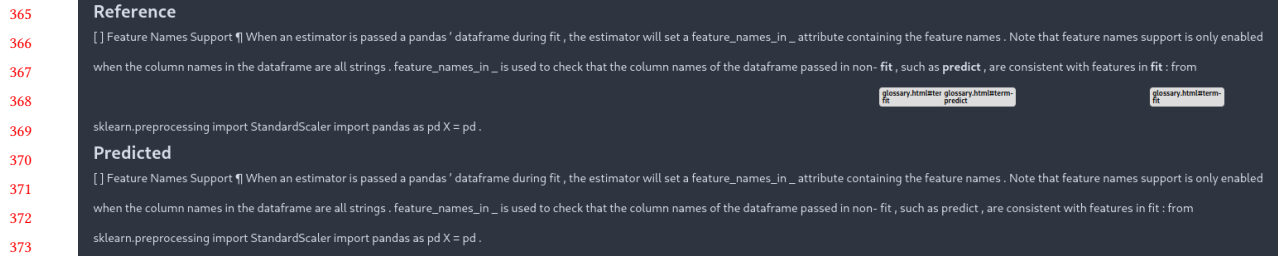


Fig. 2. An example from the testing set and the model's (lack of) predictions

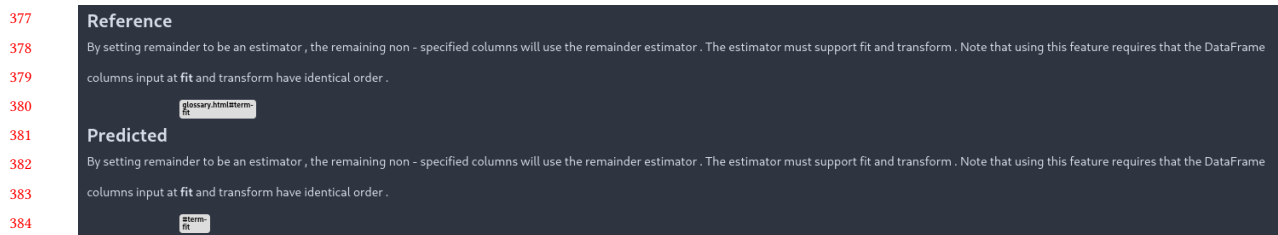


Fig. 3. Another example from the testing set and the model's prediction

terms. Furthermore, the dataset included repetitive automatically generated text, which is essentially many instances of a template filled in with particular values. This has the effect of the model learning very specific contexts for links appearing in these templated texts. These biases in the dataset account for many of the instances where the model achieved perfect scores. In future iterations of the system, it would certainly be preferable to include an option to filter automatically generated text, as this not only biases the model but also makes the model learn things that will likely not be used, as the text is automatically generated and the system is for manually written text. A more sophisticated duplicate detection algorithm could also catch these instances of templated text.

One thing to note regarding the classification scores is that they aren't perfect, because the evaluation samples and the dataset in general is imperfect: by nature, the dataset contains false negatives (spans that could be linked but aren't). For example, there are cases where the model failed to classify the expected span of text, however it classified another equally viable span of text correctly in the same test example, but this is counted as a total failure. This presence of false negatives is a general obstacle to be overcome by the system, and may be addressed by adjusting the weight given to negative samples, or providing a means of manual annotation of the dataset by the authors.

One example of the model's hyper-specialization is for linked names appearing in the changelogs, which almost exclusively appear at the end of a line preceded by the same structure of text. Figure 4 shows some of the training examples for one such class, and figure 5 shows some predictions of the model on text that I wrote. This class had a perfect F1-score. The model appears to have learned a very specific representation requiring the name to appear at the end or at the start of the text, with a preference for it to be preceded by particular text (the words "by" and "and"). However, some useful properties of the model are demonstrated here, in particular the ability to identify similar forms such as when a middle initial is included in the name, or it is misspelled.

The "#term-fit" class is an example of an imperfectly classified class (0.6 F1-score) with a more diverse set of contexts where it appears. Figure 6 shows a sample of the training examples for this class, and 7 again shows some test cases.



Train doc data/html/processed/whats\_new/v0.22.txt  
 A FutureWarning is raised now , and this will raise an error in 0.24 . If the number of features is n't consistent and negative indexing is used , an error is raised . # 14544 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 # 15257 by Mart Willocx . - Enhancement decomposition.dict\_learning and decomposition.dict\_learning\_online now accept method\_max\_iter and pass it to decomposition.sparse\_encode . # 12650 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 - Fix ensemble . HistGradientBoostingClassifier now raises an error if categorical\_crossentropy loss is given for a binary classification problem . # 14869 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 HistGradientBoostingClassifier and ensemble . HistGradientBoostingRegressor instead . # 14907 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 # 13806 by Anaël Beaugnon . - Fix The fit in FeatureUnion now accepts fit\_params to pass to the underlying transformers . # 15119 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 NuSVC now accept a break\_ties parameter . This parameter results in predict breaking the ties according to the confidence values of decision\_function , if decision\_function\_shape='ovr' , and the number of target classes > 2 .  
 # 12557 by Adrin Jalali .  
<https://github.com/adrinjalali>

Train doc data/html/processed/whats\_new/v0.22.txt  
 - API Change The classes \_ and n\_classes \_ attributes of tree . DecisionTreeRegressor are now deprecated . # 15028 by Mei Guan , Nicolas Hug , and Adrin Jalali .  
<https://github.com/nicolashug> <https://github.com/adrinjalali>

Fig. 4. Training examples for the “<https://github.com/adrinjalali>” class

Adrin Jalali  
<https://github.com/adrinjalali>

jasdvnjlksnweoiurhincvjdasklfhu Adrin Jalali asjdfnljksdnjlkwfnjwelfnjklksaf  
 Jalali

Adrin

Adrin F. Jalali did something  
<https://github.com/adrinjalali>

Jalali Adrin

A bunch of nonsense Adrin Jalali never seen before in the docs  
 A bunch of nonsense Adrin Jalali has never seen before in the docs  
 Adrin Jalali did something again in sklearn  
 Adrin Jalali did something again in sklearn by Adrin Jalali  
<https://github.com/adrinjalali>

Adrin Jalali and Adrin Jalali  
<https://github.com/adrinjalali>

Adrin Jlali  
<https://github.com/adrinjalali>

Fig. 5. Predictions for the “<https://github.com/adrinjalali>” class

This example demonstrates the model’s utilization of context to learn more sophisticated representations for the classes. In particular, we see that the model correctly does not classify the usage of “fit” as a verb in the context of clothing as an instance of this class. However we also see some potential over-specialization, where the model seems to be requiring the word “estimator” to appear in some context before it will classify “fit”. However this requirement is nuanced as demonstrated by the last example, where the presence “estimator” is not sufficient to classify “fit” because the rest of the context doesn’t match. It should also be noted that this example is from an earlier version of the model, and it is just being used to demonstrate the model’s context learning capabilities. In a later version of the model trained on more balanced training data, it seems to have reduced this over-specialization.

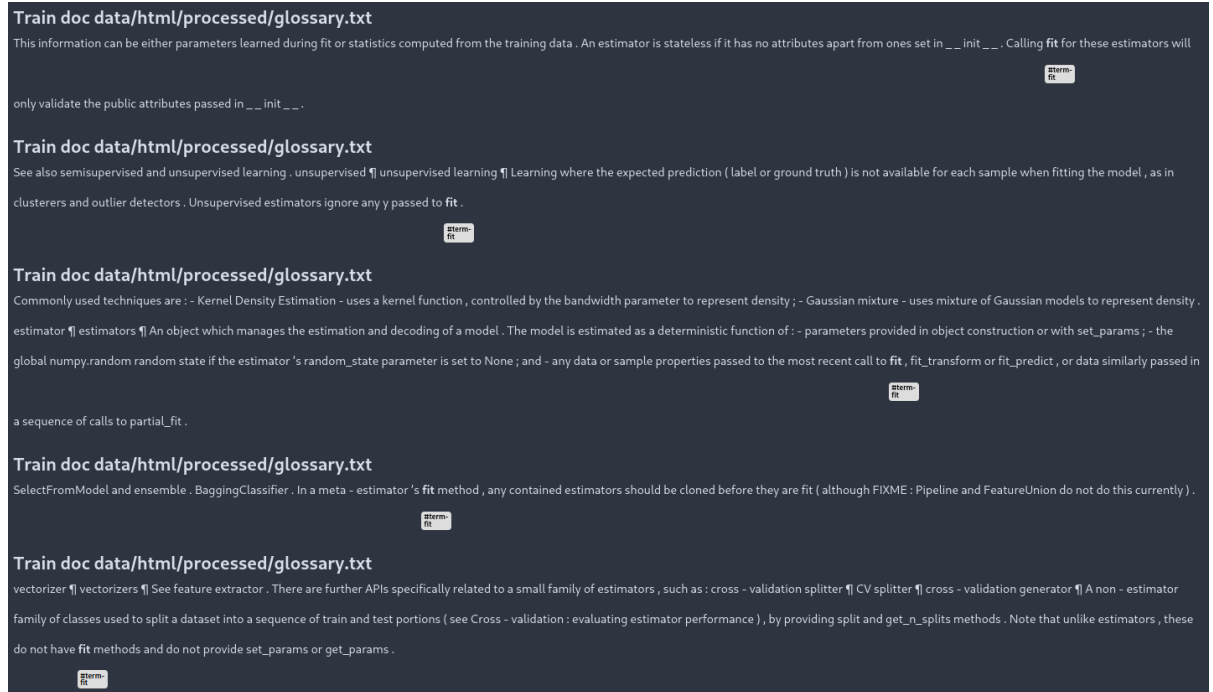


Fig. 6. Training examples for the “#term-fit” class

In conclusion, the system demonstrates potential to become a sophisticated and “intelligent” solution for corpus-tuned automatic hyperlinking, however it requires more work to improve the system’s data processing to reduce bias in the data, and the system may require some manual intervention by corpus authors to produce the best results.

Future work would center around improving the system, likely focusing on the following avenues:

- Handling false negatives: much of the text by nature contains false negatives. If these false negatives get into the training data it could give the model the wrong idea.
  - One potential improvement may be to reduce the context around examples, which would reduce the likelihood of false negatives appearing in training
  - Could explore an approach like: <https://github.com/doccano/spacy-partial-tagger>
- Reducing bias in the data: data is very biased towards unique names and biased against terms.

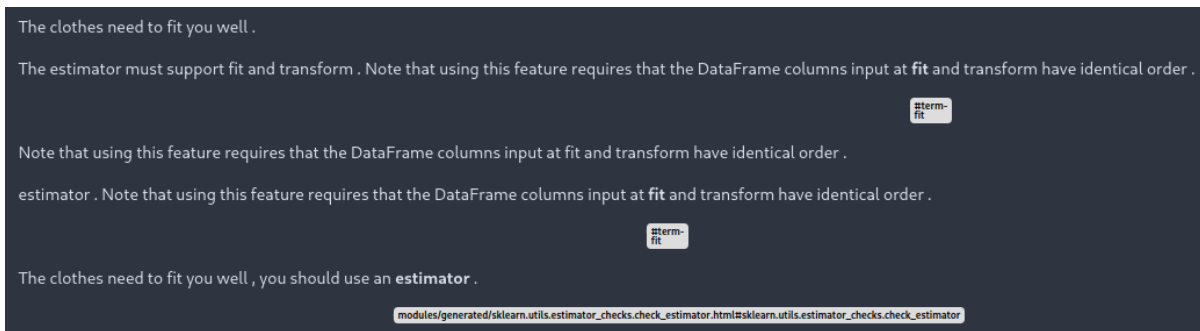


Fig. 7. Predictions for the “#term-fit” class

- This may just be a property of the dataset, in which case manual annotation by the corpus authors would be required before the system could improve
- However this may be improved by providing document filter options, for example to ignore things like API documentation, changelogs, and generated text
- Some potential data pre-processing improvements:
  - Some equal links are not caught by normalization, in particular links to a particular part of a page, from within the page and without, for example “#pipeline” and “path/to/pipeline.html#pipeline”. This could be handled by incorporating the file path of the document in which the link appears during normalization.
  - Try to increase diversity of training samples by prioritizing unique span text during undersampling. The current undersampling technique effectively chooses random samples, however an algorithm which encourage more diversity may improve the system.
  - imbalanced-learn didn’t work well for this task because it doesn’t support multi-label classification, however if some of the imbalanced learning methods could be adapted to a multi-label problem they may offer some improvement to the system.
  - Augmentation with something like [augmenty](#) [2] may provide improvement.
  - One potential augmentation would be to vary the amount of context in an example, to reduce the potential of the model leaning too much on a particular amount of context.
  - Explore minimum alpha-numeric character requirement of 1 instead of 3: the current pre-processing requires link text to have at least 3 alpha-numeric characters, however some spans like “X” and “Xt” are being removed which may provide useful data.
  - The sentence sliding window technique could be improved by ensuring the given span is in the middle of the set of sentences. The current technique allows it to be in the end sentences, making for inconsistent contexts in some cases.
  - More experimentation could be performed with the minimum class occurrence threshold: a threshold of 15 was used for most experiments, however a lower threshold may still allow for a useful model while increasing the number of classes.
- <https://spacy.io/universe/project/skweak> may be useful for ensembling, for example combining the machine learning pipeline with rule-based classifiers

- A tool such as <https://spacy.io/universe/project/spacyfishing> can be integrated into the system to facilitate pre-built wikipedia linking on top of the corpus specific links

## REFERENCES

- [1] Silviu Cucerzan. 2007. Large-Scale Named Entity Disambiguation Based on Wikipedia Data. In *Proceedings of the 2007 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning (EMNLP-CoNLL)*. Association for Computational Linguistics, Prague, Czech Republic, 708–716. <https://aclanthology.org/D07-1074>
- [2] Kenneth Enevoldsen. 2023. *Augmenty: The cherry on top of your NLP pipeline*. <https://github.com/KennethEnevoldsen/augmenty>
- [3] Hugging Face. 2023. *Hugging Face*. <https://huggingface.co/>
- [4] Ines; Van Landeghem Sofie; Boyd Adriane Honnibal, Matthew; Montani. 2020. spaCy: Industrial-strength Natural Language Processing in Python. (2020).
- [5] imbalanced learn. 2023. *imbalanced-learn*. <https://imbalanced-learn.org/stable/index.html>
- [6] Jessica Lin and Amir Zeldes. 2021. WikiGUM: Exhaustive Entity Linking for Wikification in 12 Genres. *CoRR* abs/2109.07449 (2021). arXiv:2109.07449 <https://arxiv.org/abs/2109.07449>
- [7] Rada Mihalcea and Andras Csomai. 2007. Wikify!: Linking Documents to Encyclopedic Knowledge. In *Proceedings of the Sixteenth ACM Conference on Conference on Information and Knowledge Management (Lisbon, Portugal) (CIKM '07)*. ACM, New York, NY, USA, 233–242. <https://doi.org/10.1145/1321440.1321475>
- [8] Pandoc. 2023. *Pandoc, a universal document converter*. <https://pandoc.org/>
- [9] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. 2011. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* 12 (2011), 2825–2830.
- [10] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *ArXiv* abs/1910.01108 (2019).
- [11] Scikit-learn. 2023. 3.3. *Metrics and scoring: quantifying the quality of predictions - Classification metrics*. [https://scikit-learn.org/stable/modules/model\\_evaluation.html#classification-metrics](https://scikit-learn.org/stable/modules/model_evaluation.html#classification-metrics)
- [12] Scikit-learn. 2023. *Glossary of Common Terms and API Elements - Feature*. <https://scikit-learn.org/stable/glossary.html#term-feature>
- [13] Scikit-learn. 2023. *Glossary of Common Terms and API Elements - X*. <https://scikit-learn.org/stable/glossary.html#term-X>
- [14] Scikit-learn. 2023. *Glossary of Common Terms and API Elements - Y*. <https://scikit-learn.org/stable/glossary.html#term-Y>
- [15] Konstantinos Sechidis, Grigorios Tsoumakas, and Ioannis Vlahavas. 2011. On the stratification of multi-label data. *Machine Learning and Knowledge Discovery in Databases* (2011), 145–158.
- [16] Haihao Shen, Ofir Zafrir, Bo Dong, Hengyu Meng, Xinyu Ye, Zhe Wang, Yi Ding, Hanwen Chang, Guy Boudoukh, and Moshe Wasserblat. 2022. Fast DistilBERT on CPUs. *arXiv*:2211.07715 [cs.CL]
- [17] P. Szymański and T. Kajdanowicz. 2017. A scikit-based Python environment for performing multi-label classification. *ArXiv e-prints* (Feb. 2017). *arXiv*:1702.01460 [cs.LG]
- [18] Piotr Szymański and Tomasz Kajdanowicz. 2017. A Network Perspective on Stratification of Multi-Label Data. In *Proceedings of the First International Workshop on Learning with Imbalanced Domains: Theory and Applications (Proceedings of Machine Learning Research, Vol. 74)*, Luís Torgo, Bartosz Krawczyk, Paula Branco, and Nuno Moniz (Eds.). PMLR, ECML-PKDD, Skopje, Macedonia, 22–35.
- [19] Wikipedia. 2023. *Entity linking*. [https://en.wikipedia.org/wiki/Entity\\_linking](https://en.wikipedia.org/wiki/Entity_linking)
- [20] Xiaoyao Yin, Yangchen Huang, Bin Zhou, Aiping Li, Long Lan, and Yan Jia. 2019. Deep Entity Linking via Eliminating Semantic Ambiguity With BERT. *IEEE Access* 7 (2019), 169434–169445. <https://doi.org/10.1109/ACCESS.2019.2955498>